

```
In [3]: pip install yellowbrick
```

```
Collecting yellowbrick
  Downloading yellowbrick-1.5-py3-none-any.whl (282 kB)
    ----- 282.6/282.6 kB 758.7 kB/s eta 0:00:00
Requirement already satisfied: scikit-learn>=1.0.0 in c:\users\vaibh\anaconda3\lib\site-packages (from yellowbrick) (1.0.2)
Requirement already satisfied: cycycler>=0.10.0 in c:\users\vaibh\anaconda3\lib\site-packages (from yellowbrick) (0.11.0)
Requirement already satisfied: numpy>=1.16.0 in c:\users\vaibh\anaconda3\lib\site-packages (from yellowbrick) (1.21.5)
Requirement already satisfied: matplotlib!=3.0.0,>=2.0.2 in c:\users\vaibh\anaconda3\lib\site-packages (from yellowbrick) (3.5.2)
Requirement already satisfied: scipy>=1.0.0 in c:\users\vaibh\anaconda3\lib\site-packages (from yellowbrick) (1.9.1)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\vaibh\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (4.25.0)
Requirement already satisfied: pyparsing>=2.2.1 in c:\users\vaibh\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\vaibh\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (2.8.2)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\vaibh\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (1.4.2)
Requirement already satisfied: packaging>=20.0 in c:\users\vaibh\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (21.3)
Requirement already satisfied: pillow>=6.2.0 in c:\users\vaibh\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (9.2.0)
Requirement already satisfied: joblib>=0.11 in c:\users\vaibh\anaconda3\lib\site-packages (from scikit-learn>=1.0.0->yellowbrick) (1.1.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\vaibh\anaconda3\lib\site-packages (from scikit-learn>=1.0.0->yellowbrick) (2.2.0)
Requirement already satisfied: six>=1.5 in c:\users\vaibh\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.0.0,>=2.0.2->yellowbrick) (1.16.0)
Installing collected packages: yellowbrick
Successfully installed yellowbrick-1.5
Note: you may need to restart the kernel to use updated packages.
```

```
In [6]: import warnings
with warnings.catch_warnings():
    warnings.simplefilter(action='ignore', category=FutureWarning)
```

```
In [4]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LassoCV
from sklearn.metrics import mean_squared_error, r2_score
```

```

import matplotlib.pyplot as plt
from sklearn.model_selection import ShuffleSplit
from yellowbrick.regressor import PredictionError
from yellowbrick.model_selection import LearningCurve
from sklearn.decomposition import PCA
import seaborn as sns

np.random.seed(1)

class DataScienceProject:
    def __init__(self):
        pass

    def load_data(self, file_path):
        # Load data using pandas
        data = pd.read_csv(file_path)
        return data

    def report_missing_values(self, df ):
        # Calculate the number of missing values per column
        missing_values = df.isnull().sum()
        missing_report = pd.DataFrame(missing_values, columns=['missing_values'])
        missing_report = missing_report[missing_report['missing_values'] > 0]

        # Suggest imputation values
        imputation_values = {}
        for column in missing_report.index:
            if df[column].dtype in ['int64', 'float64']:
                skewness = df[column].skew()
                if abs(skewness) > 0.5:
                    imputation_value = df[column].median()
                    imputation_values[column] = ('median', imputation_value)
                else:
                    imputation_value = df[column].mean()
                    imputation_values[column] = ('mean', imputation_value)
            else:
                imputation_value = df[column].mode()[0]
                imputation_values[column] = ('mode', imputation_value)

        return imputation_values

    def apply_imputations(self, df, imputation_values):
        for column, (strategy, value) in imputation_values.items():
            if strategy in ['mean', 'median', 'mode']:

```

```

        df[column].fillna(value, inplace=True)
    # print("after imp", df.isna().sum())
    return df

def visualize_data(self, data, title_suffix=''):
    """
    Visualizes distributions of numerical and categorical features in the dataset.

    Args:
    data (DataFrame): The dataset to visualize.
    title_suffix (str): A suffix for the plot title to distinguish between original and preprocessed data.
    """
    numerical_cols = data.select_dtypes(include=['int64', 'float64']).columns
    categorical_cols = data.select_dtypes(include=['object']).columns

    # Plot for numerical features
    for col in numerical_cols:
        plt.figure(figsize=(8, 4))
        sns.histplot(data[col], kde=True)
        plt.title(f'Distribution of {col} {title_suffix}')
        plt.xlabel(col)
        plt.ylabel('Frequency')
        plt.show()

    # Plot for categorical features
    for col in categorical_cols:
        plt.figure(figsize=(8, 4))
        sns.countplot(x=col, data=data)
        plt.title(f'Distribution of {col} {title_suffix}')
        plt.xlabel(col)
        plt.ylabel('Count')
        plt.show()

def select_data_within_iqr(self, df, iqr_factor=1.5):
    Q1 = df.quantile(0.2)
    Q3 = df.quantile(0.8)
    IQR = Q3 - Q1

    lower_bound = Q1 - (iqr_factor * IQR)
    upper_bound = Q3 + (iqr_factor * IQR)

    # Select only the rows where each column value is within the IQR bounds
    selected_data = df[~((df < lower_bound) | (df > upper_bound)).any(axis=1)]

```

```

        # print("after iqr", selected_data.isna().sum())
        return selected_data

def apply_log_transformation(self, df, target_column, skew_threshold=0.5):
    """
    Applies log transformation to highly skewed columns.

    Args:
    df (DataFrame): The dataframe containing the data.
    skew_threshold (float): The threshold to identify highly skewed columns.

    Returns:
    DataFrame: The dataframe with log-transformed columns.
    """
    for column in df.select_dtypes(include=['float64', 'int64']):
        if df[column].skew() > skew_threshold and column != target_column :
            df['Log_' + column] = np.log1p(df[column])

    # print(df.columns)
    # print("after log", df.isna().sum())
    return df

def preprocess_data(self, data, target_column):
    # Report and apply imputations, and handle outliers
    imputation_values = self.report_missing_values(data)
    data = self.apply_imputations(data, imputation_values)
    data = self.apply_log_transformation(data, target_column)
    data = self.select_data_within_iqr(data)
    return data

def train_model(self, training_data, target_column):
    # Modify to return X_train, X_test, y_train, y_test for Lasso analysis
    X = training_data.drop(target_column, axis=1)
    y = training_data[target_column]
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)

    model = LinearRegression()
    model.fit(X_train, y_train)
    return model, X_train, X_test, y_train, y_test

def plot_actual_vs_predicted(self, y_test, y_predict, model_type):
    """
    Plots the actual vs predicted values.

    Args:

```

```

y_test (array-like): The true values of the target variable.
y_predict (array-like): The predicted values by the model.
model_type (str): Type of the model ('Baseline' or 'Main').
"""

plt.figure(figsize=(16, 6))
plt.title(f"{model_type} Model: Actual vs Predicted Values")
x_points = list(range(len(y_test)))
plt.plot(x_points, y_test, label='Actual Values', marker='o')
plt.plot(x_points, y_predict, label='Predicted Values', marker='x')
plt.xlabel('Data Points')
plt.ylabel('Target Variable')
plt.legend()
plt.show()

```

```

def plot_learning_curve_and_prediction_error(self, model, X_train, X_test, y_train, y_test, model_type):
    """
    Plots the learning curve and prediction error using Yellowbrick.

    Args:
    model: The trained model.
    X_train, X_test, y_train, y_test: Training and testing data.
    model_type (str): Type of the model ('Baseline' or 'Main').
    """

    # Learning Curve
    plt.figure(figsize=(10, 6))
    lc_viz = LearningCurve(model, cv=5, scoring='r2', n_jobs=4, train_sizes=np.linspace(0.1, 1.0, 10))
    lc_viz.fit(X_train, y_train)
    lc_viz.set_title(f"{model_type} Model: Learning Curve")
    lc_viz.show()

    # Prediction Error
    plt.figure(figsize=(10, 6))
    pe_viz = PredictionError(model)
    pe_viz.fit(X_train, y_train)
    pe_viz.score(X_test, y_test)
    pe_viz.set_title(f"{model_type} Model: Prediction Error")
    pe_viz.show()

def feature_importance_lasso(self, X_train, y_train, X_test, y_test):
    # Create and fit the LassoCV model
    lasso = LassoCV(cv=5, random_state=0)
    lasso.fit(X_train, y_train)

```

```

best_alpha = lasso.alpha_
lasso_coef = lasso.coef_
y_pred = lasso.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Plotting feature importances
plt.figure(figsize=(10, 6))
feature_importance = pd.Series(lasso_coef, index=X_train.columns).sort_values()
feature_importance.plot(kind='barh')
plt.title('Feature Importances from Lasso Model')
plt.xlabel('Importance')
plt.ylabel('Features')
plt.show()

print(f"Best alpha: {best_alpha}")
print(f"MSE: {mse}")
print(f"R-squared: {r2:.2f}")

def make_prediction(self, model, new_data):
    # Make predictions using the trained model
    prediction = model.predict(new_data)
    return prediction

def train_and_evaluate(self, X_train, X_test, y_train, y_test):
    """
    Train a linear regression model and evaluate it.
    """
    model = LinearRegression()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    return model, mse, r2

def main_pipeline(self, file_path, target_column):
    print("==== Data Science Project Pipeline =====")
    print("[Step 1] Loading and Preprocessing Data")
    data = self.load_data(file_path)

    # Visualize original data
    self.visualize_data(data, title_suffix='(Original)')

    preprocessed_data = self.preprocess_data(data, target_column)

```

```

# Visualize preprocessed data
self.visualize_data(preprocessed_data, title_suffix='(Preprocessed)')

# Baseline Model
print("\n[Step 2] Training and Evaluating Baseline Model")
mse_baseline, r2_baseline, baseline_model, X_train_baseline, X_test_baseline, y_train_baseline, y_test_baseline = self.train_and_evaluate(preprocessed_data, target_column)
y_pred_baseline = baseline_model.predict(X_test_baseline)
self.plot_actual_vs_predicted(y_test_baseline, y_pred_baseline, "Baseline")
self.plot_learning_curve_and_prediction_error(baseline_model, X_train_baseline, X_test_baseline, y_train_baseline, y_test_baseline)

# Main Model
print("\n[Step 3] Training and Evaluating Main Model")
model, X_train, X_test, y_train, y_test = self.train_model(preprocessed_data, target_column)
_, mse_main, r2_main = self.train_and_evaluate(X_train, X_test, y_train, y_test)
y_predict_main = model.predict(X_test)
self.plot_actual_vs_predicted(y_test, y_predict_main, "Main")
self.plot_learning_curve_and_prediction_error(model, X_train, X_test, y_train, y_test, "Main")

print("\n[Step 4] Performing Lasso Feature Importance Analysis")
self.feature_importance_lasso(X_train, y_train, X_test, y_test)

print("\n==== Model Comparison Results =====")
print("Baseline Model:")
print(" - MSE: {:.3f}".format(mse_baseline))
print(" - R-squared: {:.3f}".format(r2_baseline))
print("Main Model:")
print(" - MSE: {:.3f}".format(mse_main))
print(" - R-squared: {:.3f}".format(r2_main))
print("=====")

return preprocessed_data

def compare_with_baseline(self, data, target_column):
    # Updated to return the model and train/test splits
    baseline_data = data.dropna()
    X_baseline = baseline_data.drop(target_column, axis=1)
    y_baseline = baseline_data[target_column]
    X_train_baseline, X_test_baseline, y_train_baseline, y_test_baseline = train_test_split(X_baseline, y_baseline,
    baseline_model, mse_baseline, r2_baseline = self.train_and_evaluate(X_train_baseline, X_test_baseline, y_train_baseline, y_test_baseline)
    return mse_baseline, r2_baseline, baseline_model, X_train_baseline, X_test_baseline, y_train_baseline, y_test_baseline

```

```

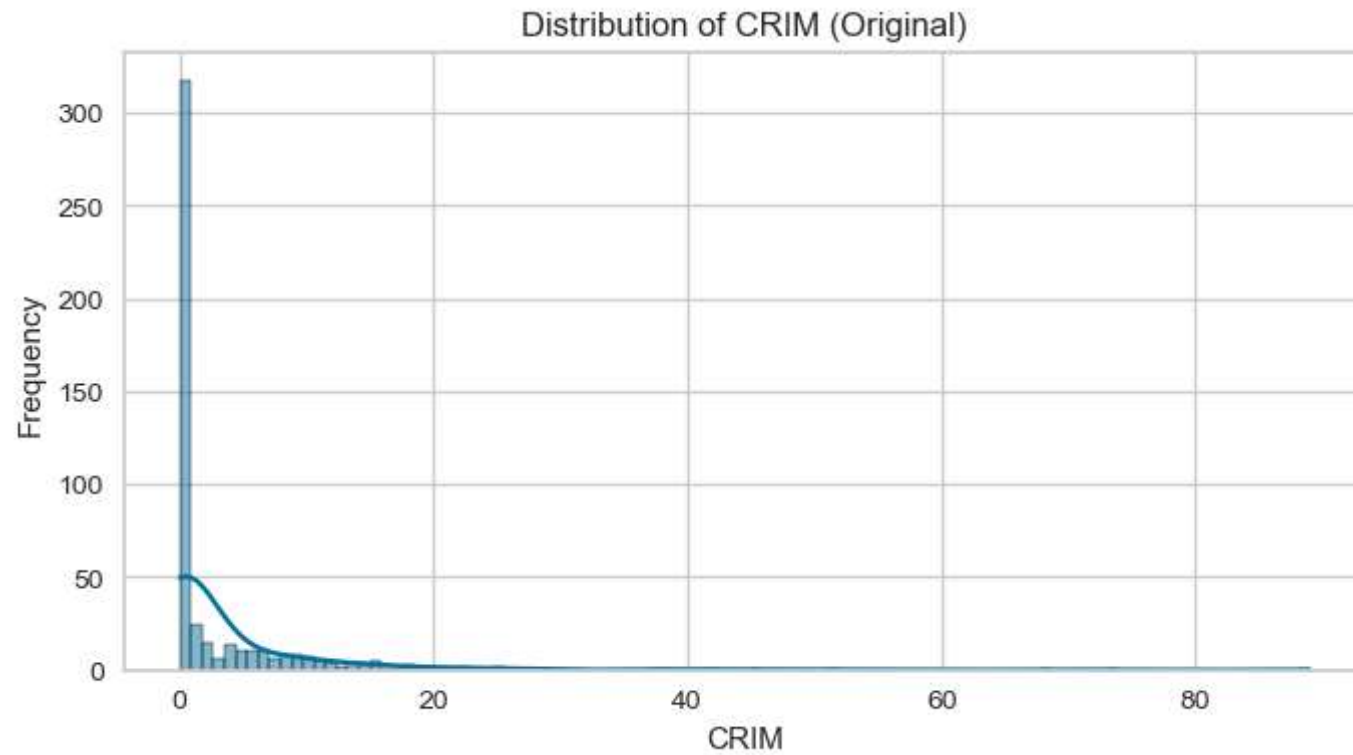
In [7]: # Create an instance of the DataScienceProject class
dsp = DataScienceProject()

```


```
# Run the main pipeline
prediction = dsp.main_pipeline("D:\Temp\HousingData.csv", 'MEDV')
```

==== Data Science Project Pipeline ====

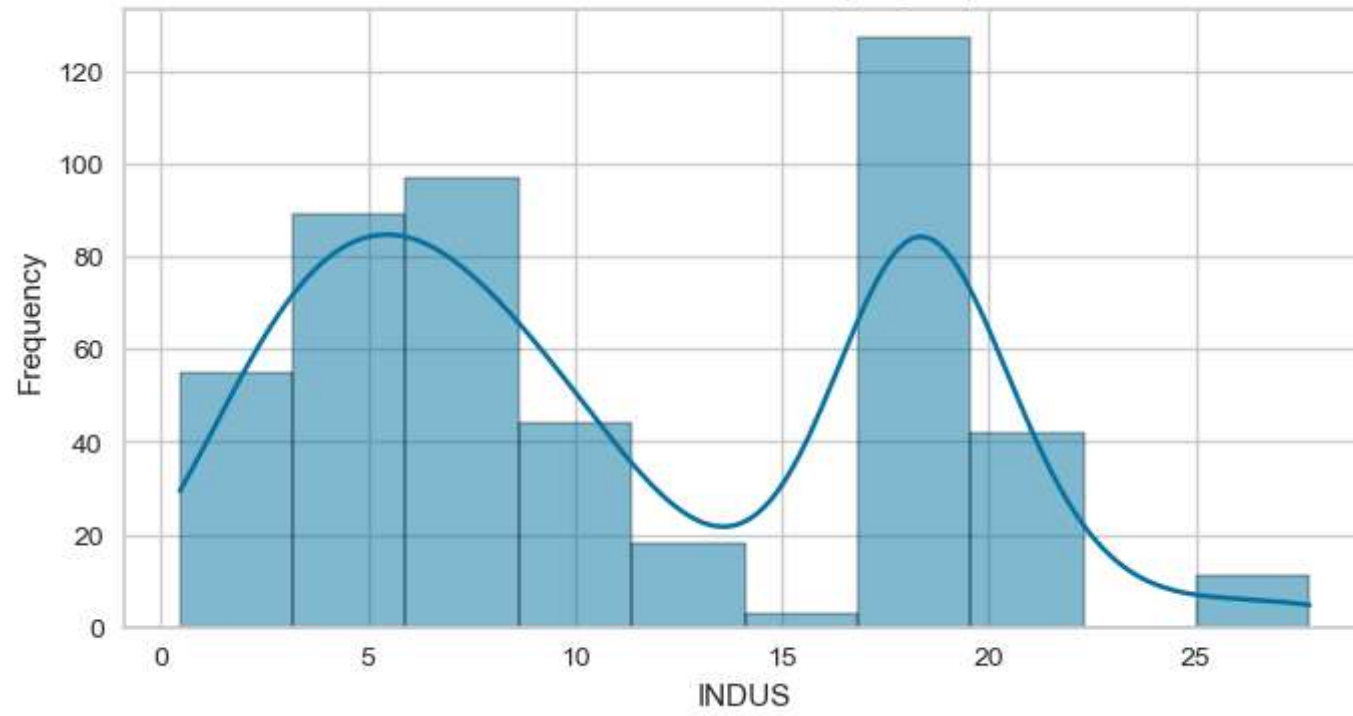
[Step 1] Loading and Preprocessing Data



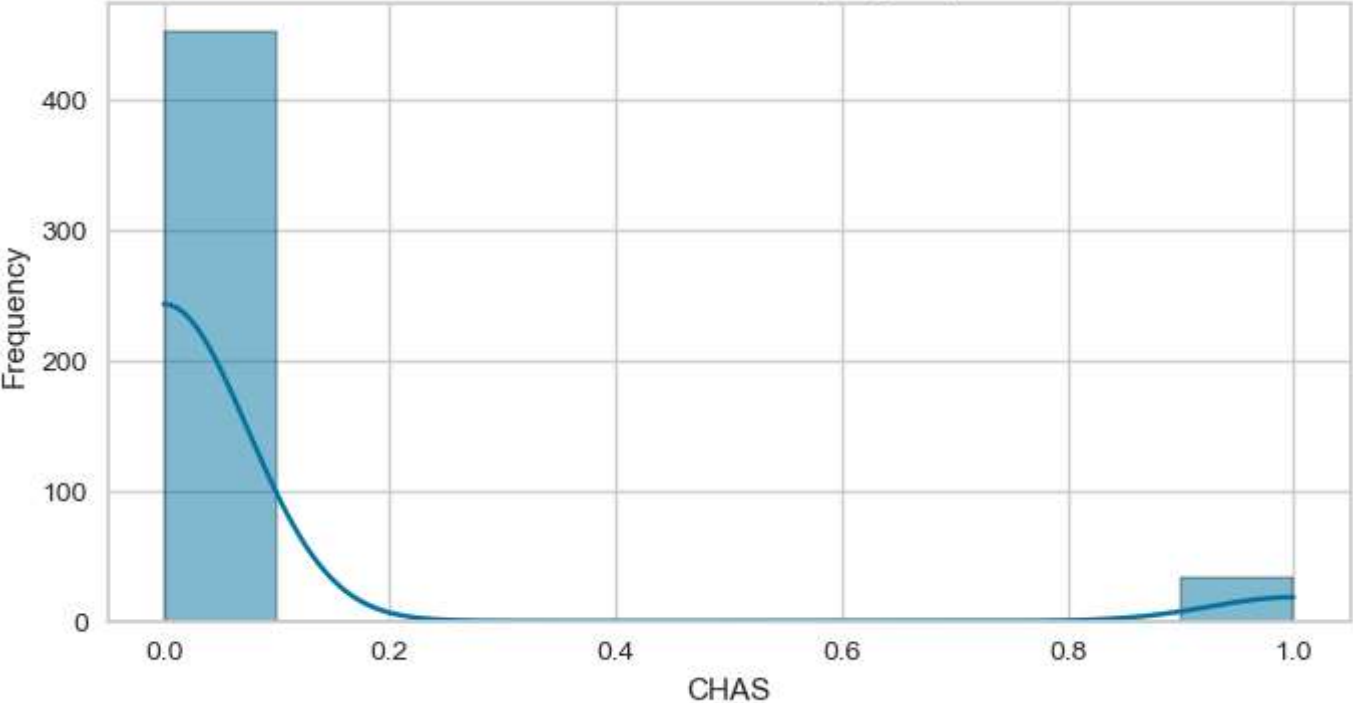
Distribution of ZN (Original)



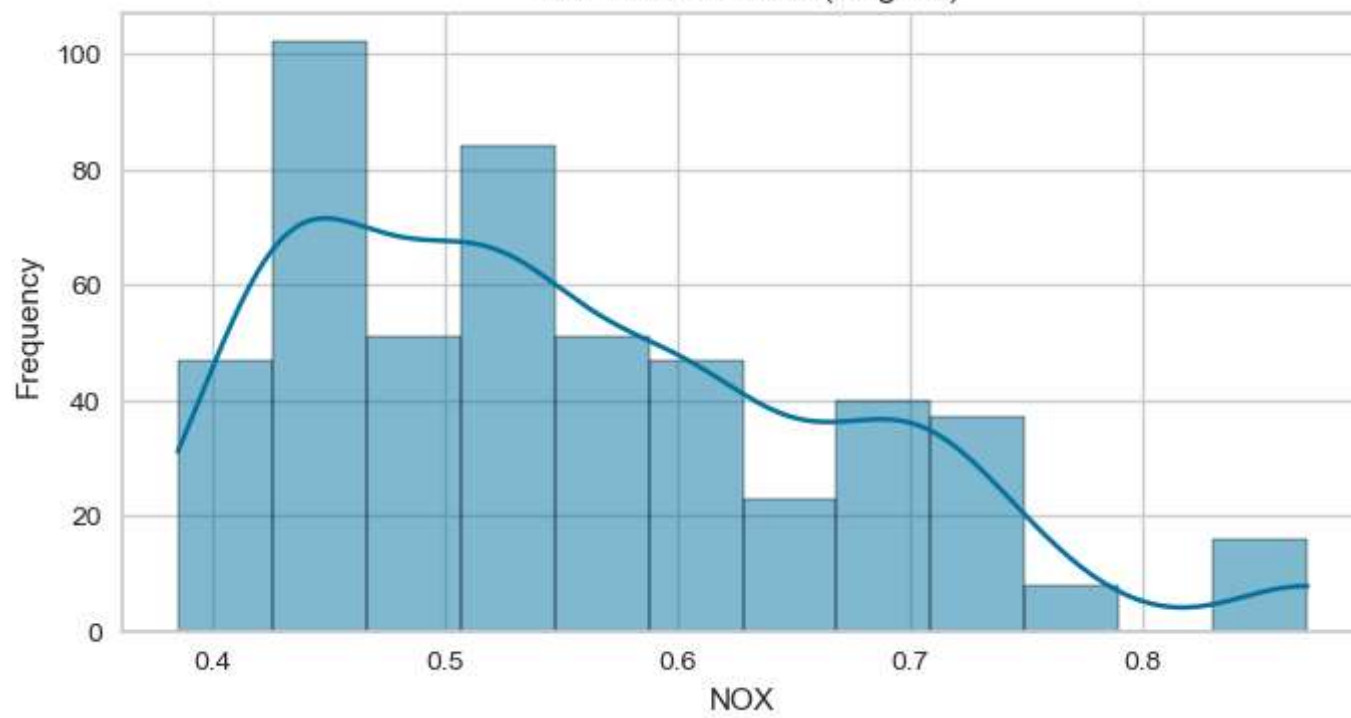
Distribution of INDUS (Original)




Distribution of CHAS (Original)




Distribution of NOX (Original)



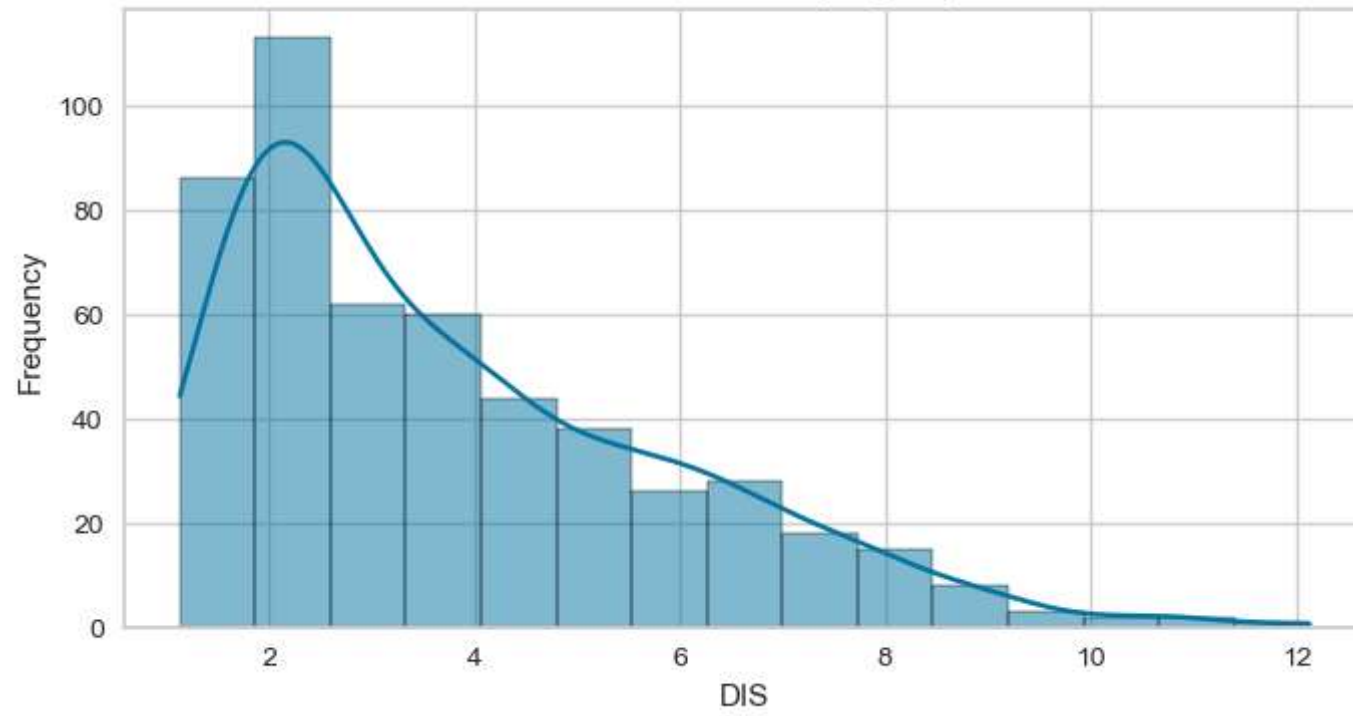
Distribution of RM (Original)



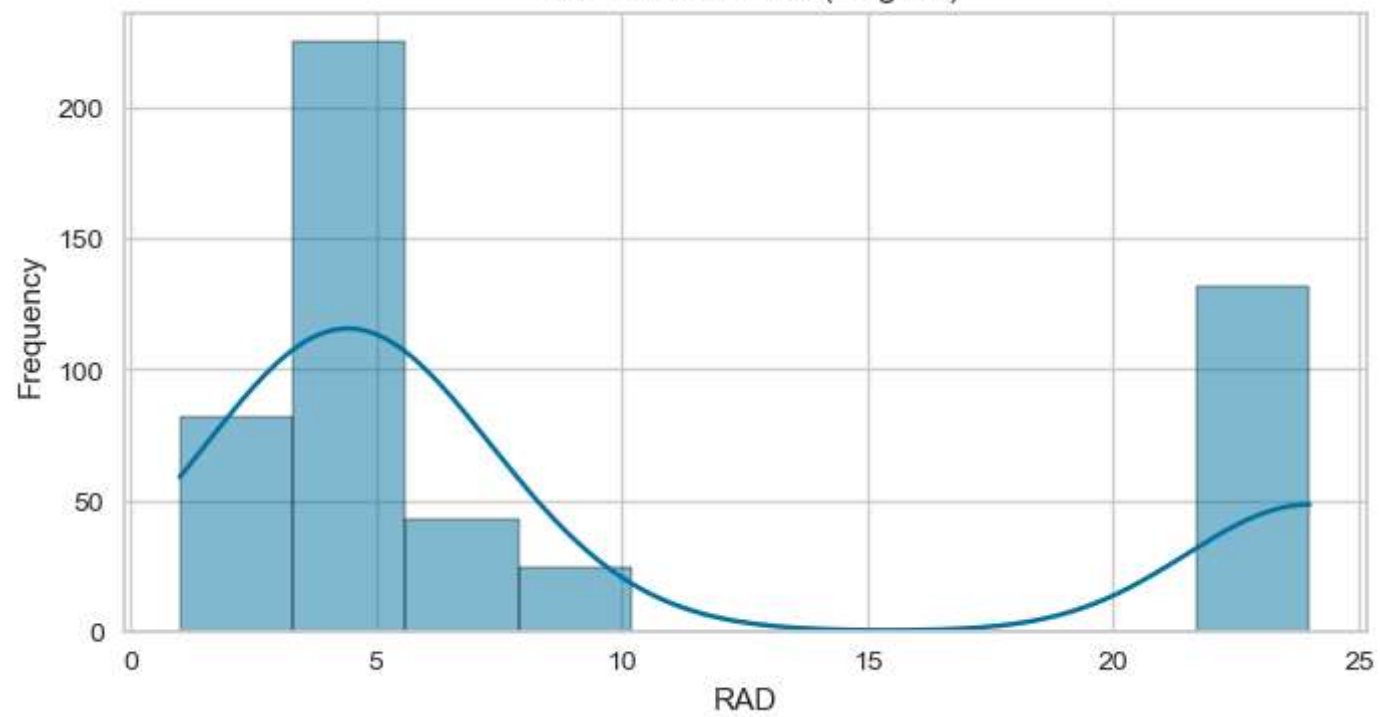
Distribution of AGE (Original)



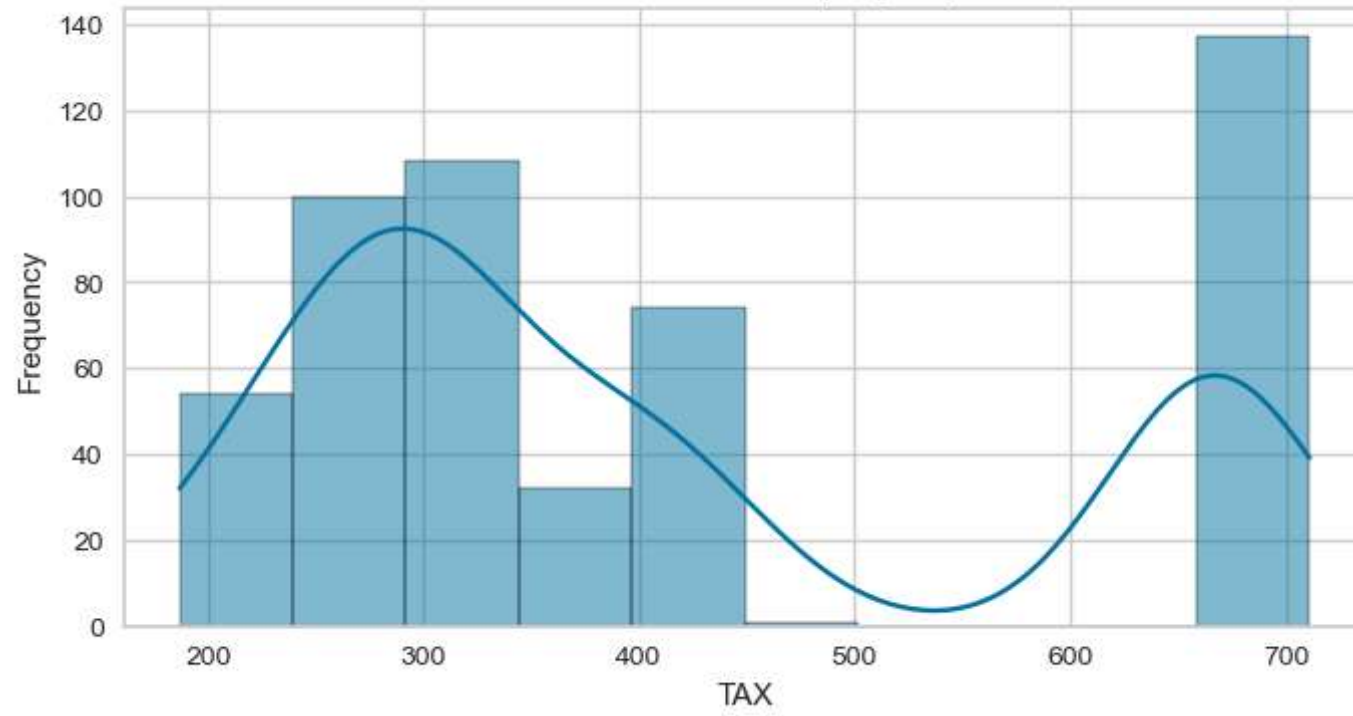
Distribution of DIS (Original)




Distribution of RAD (Original)




Distribution of TAX (Original)

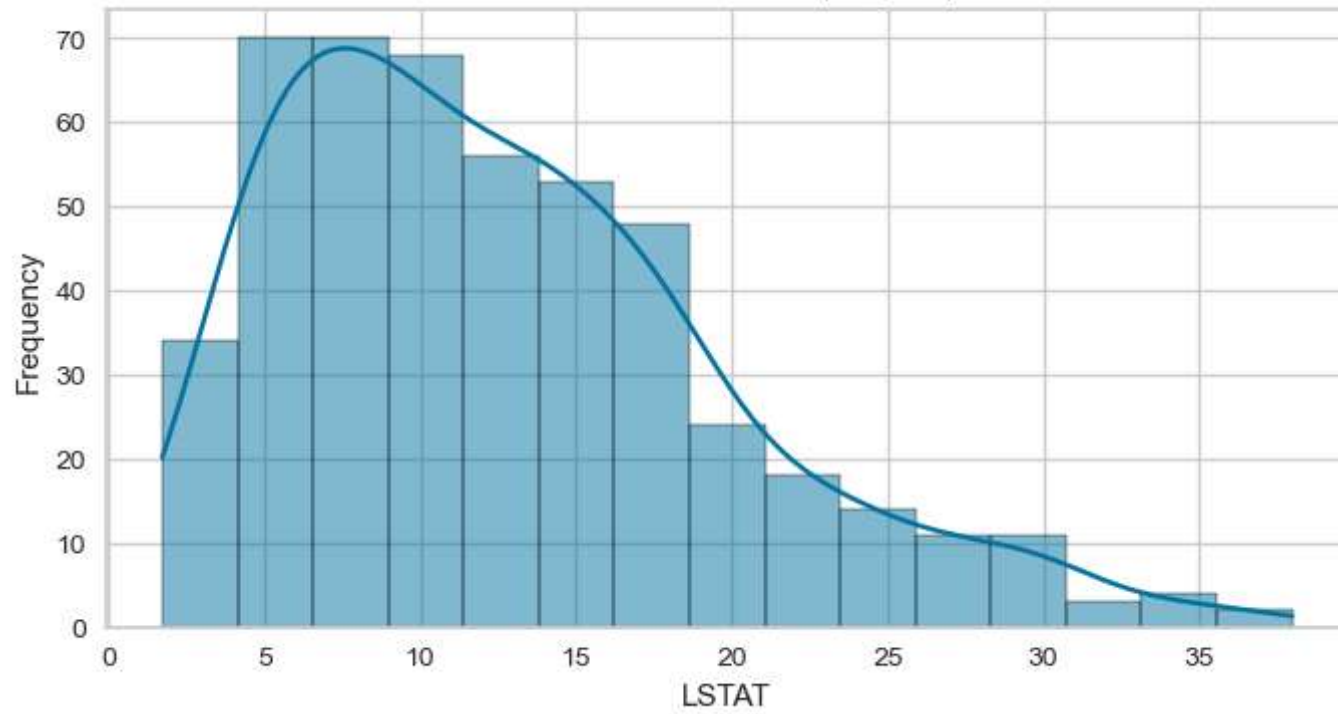



Distribution of PTRATIO (Original)



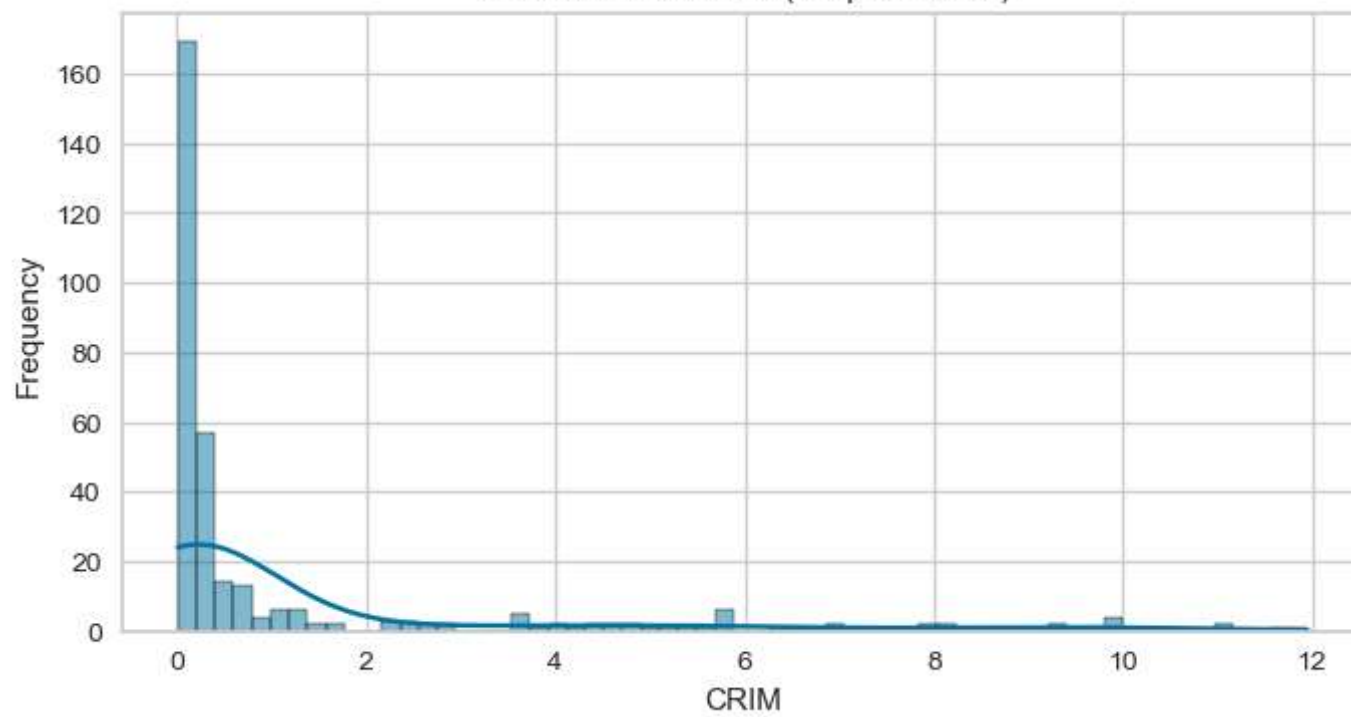


Distribution of LSTAT (Original)

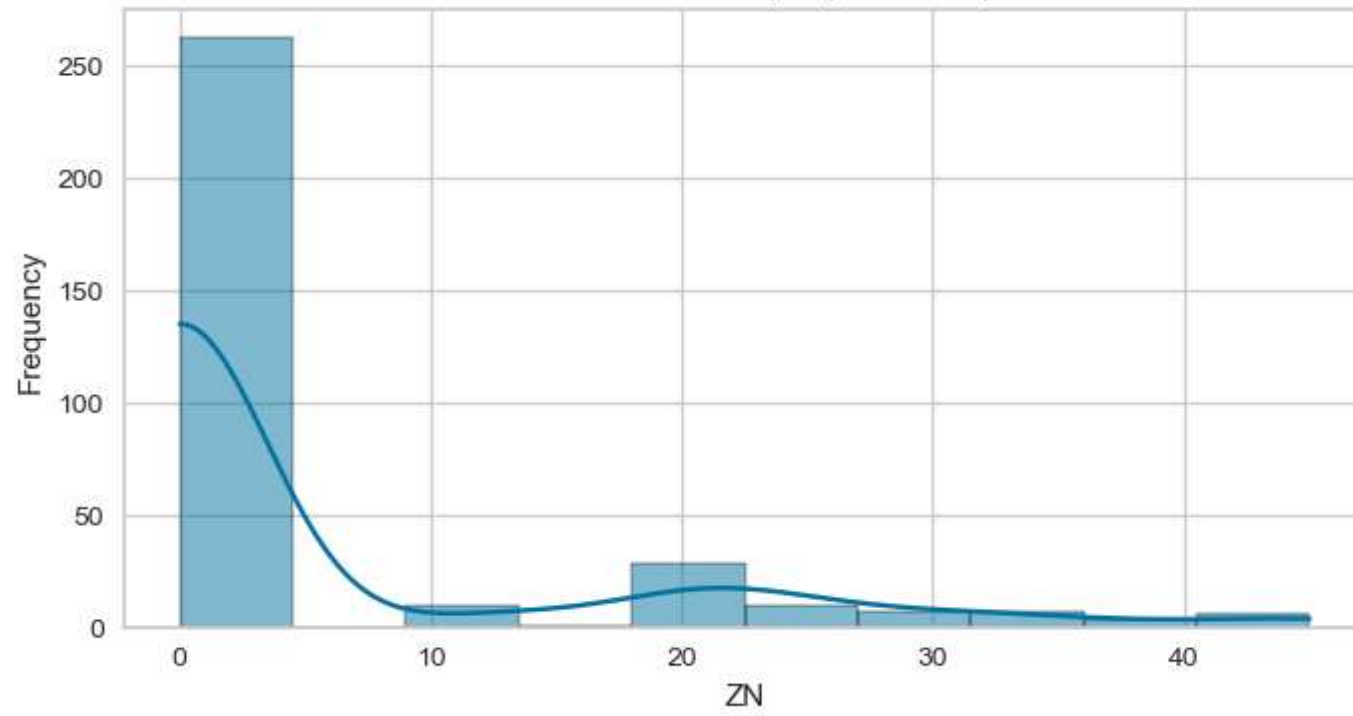


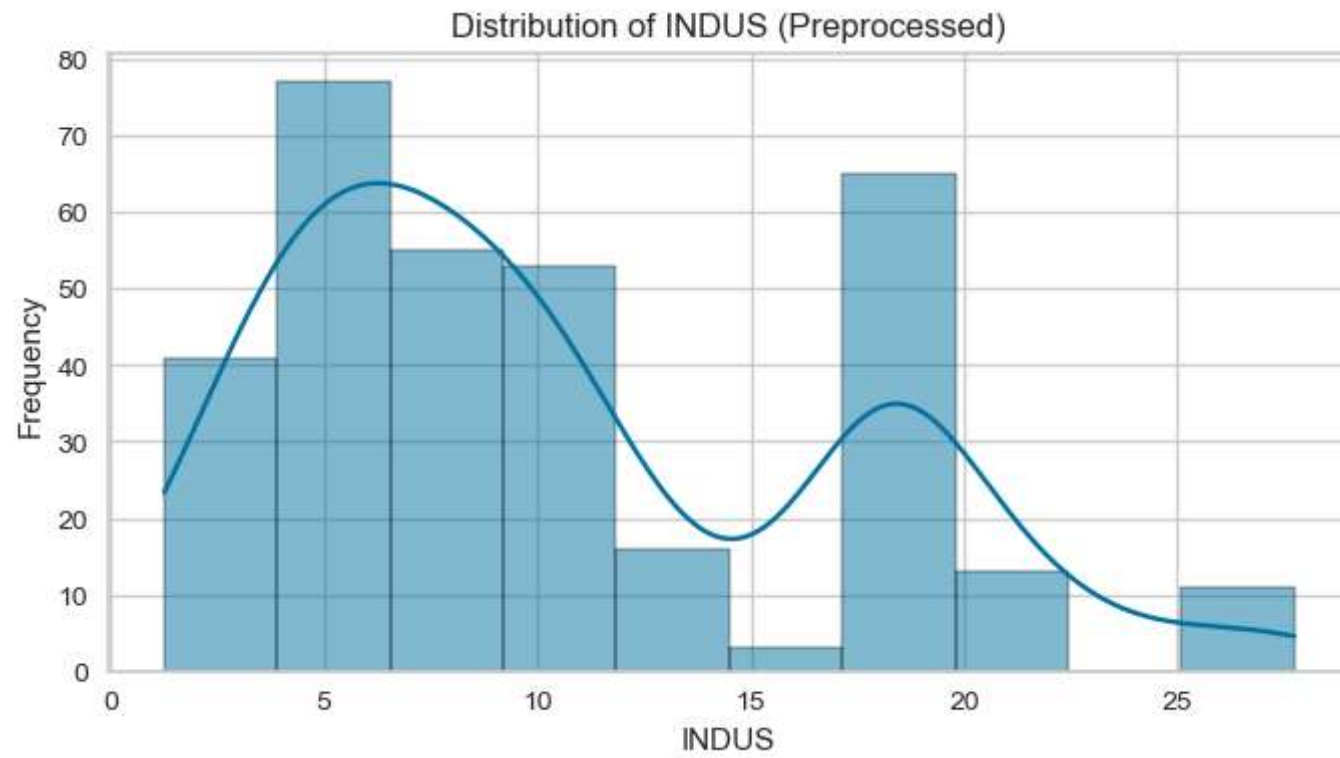


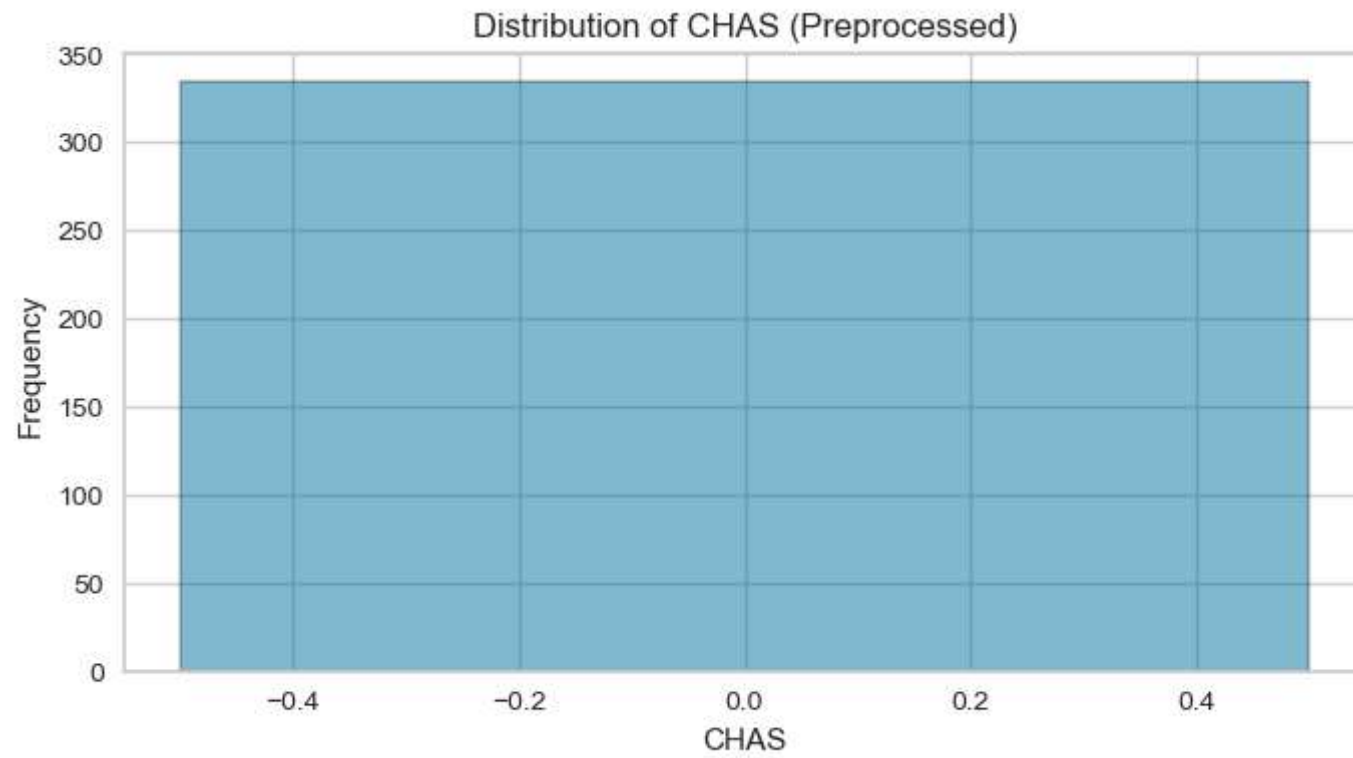
Distribution of CRIM (Preprocessed)

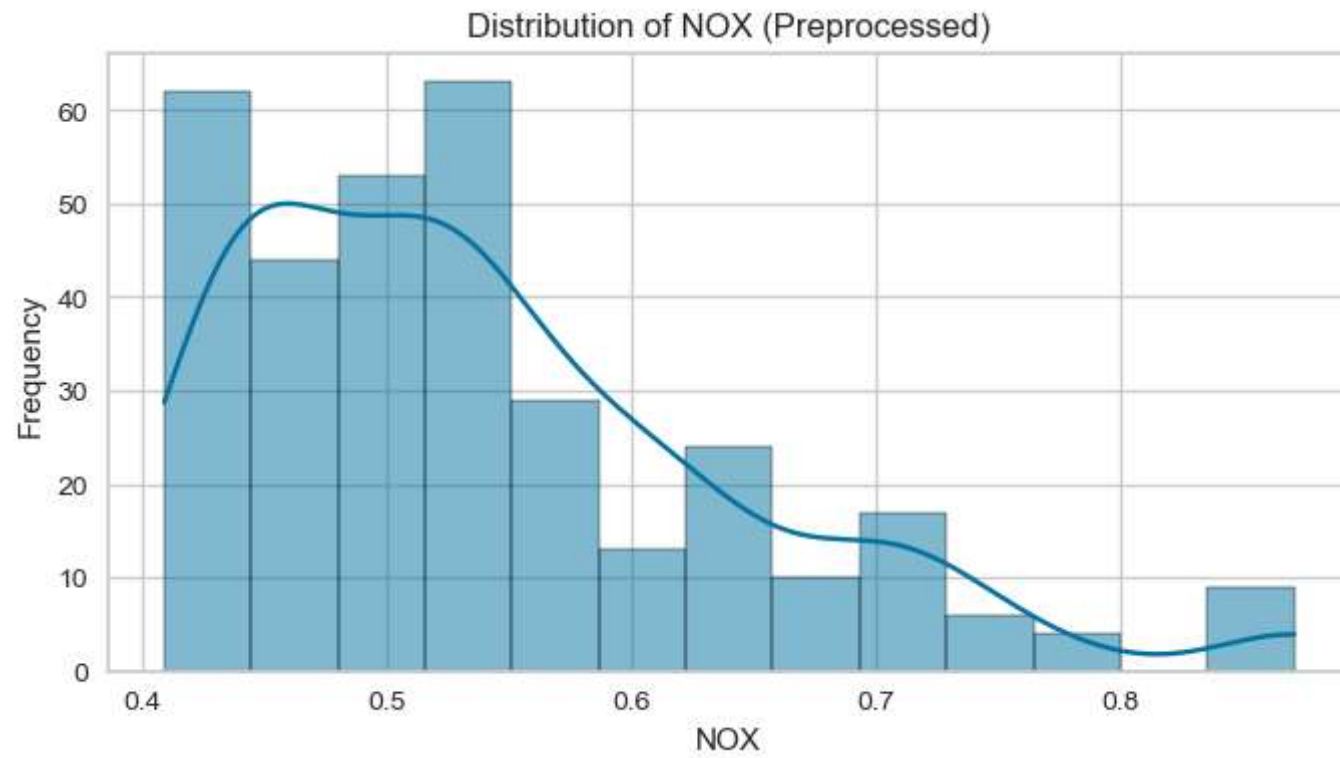


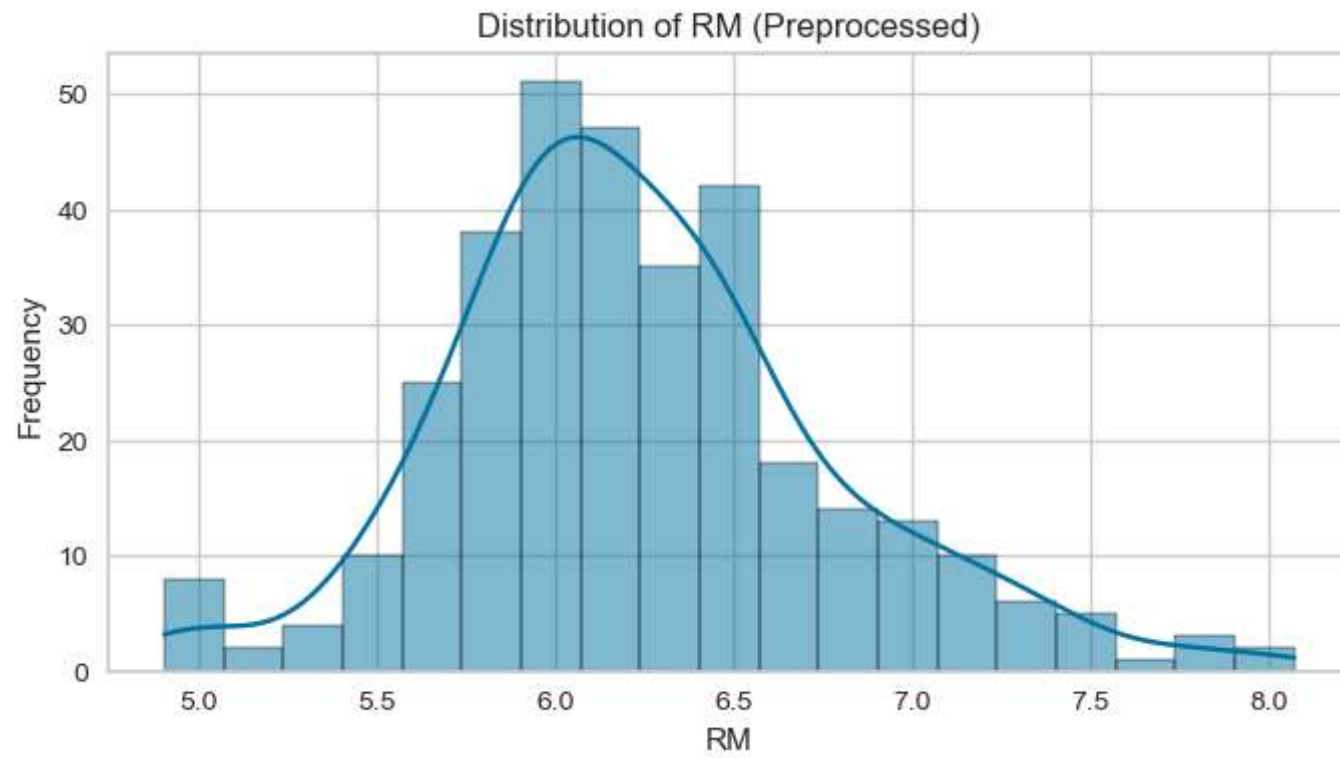
Distribution of ZN (Preprocessed)



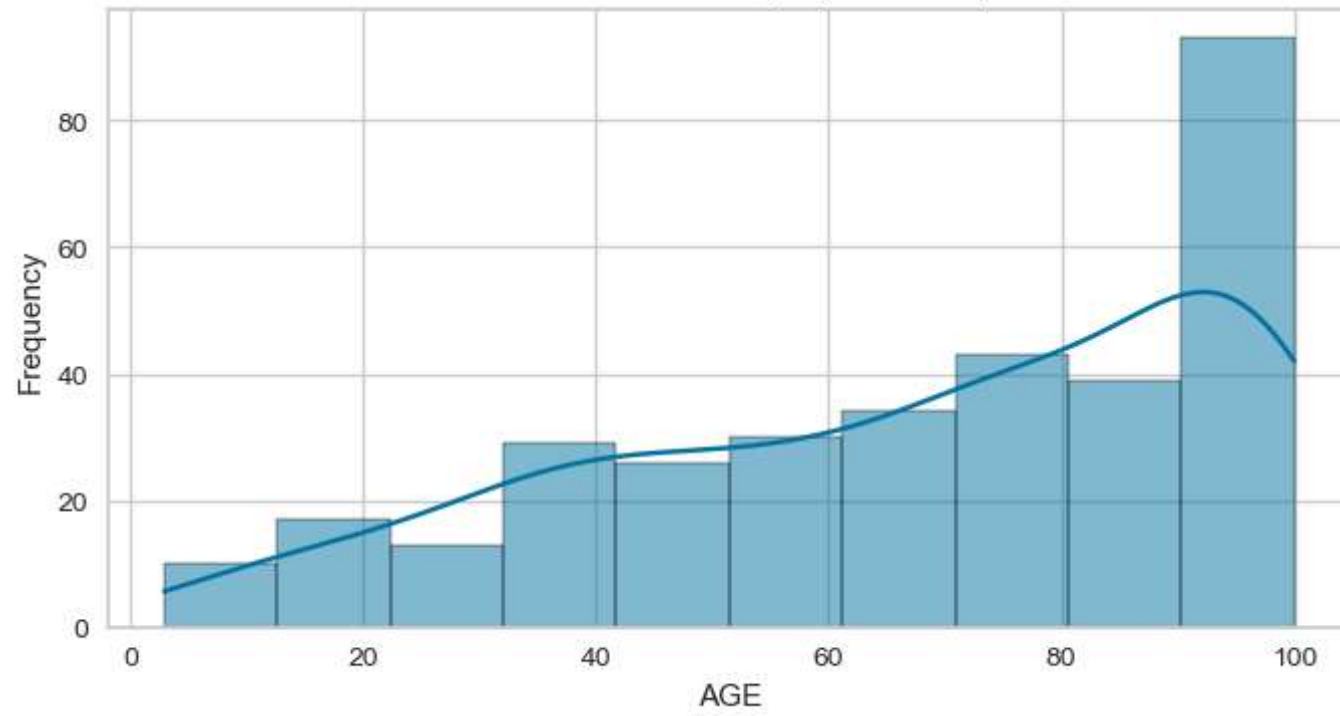










Distribution of AGE (Preprocessed)

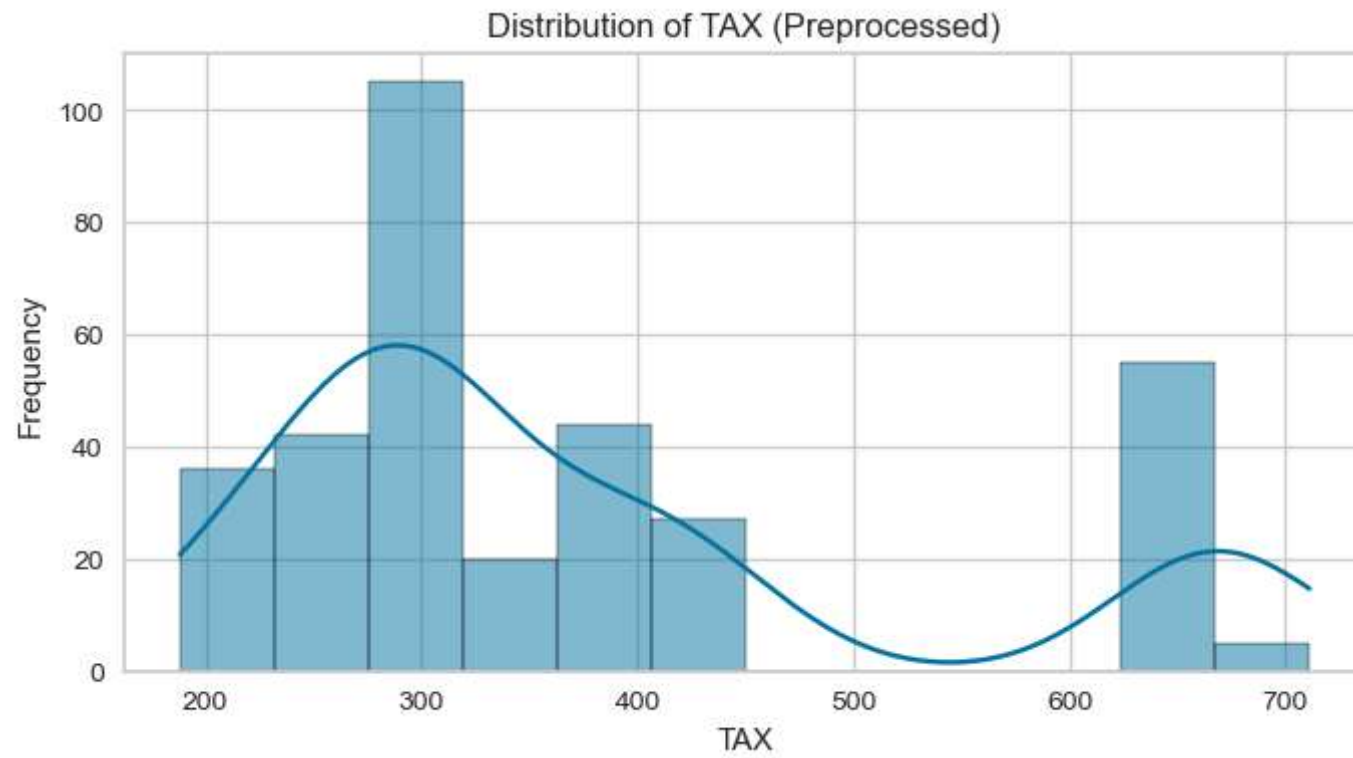


Distribution of DIS (Preprocessed)

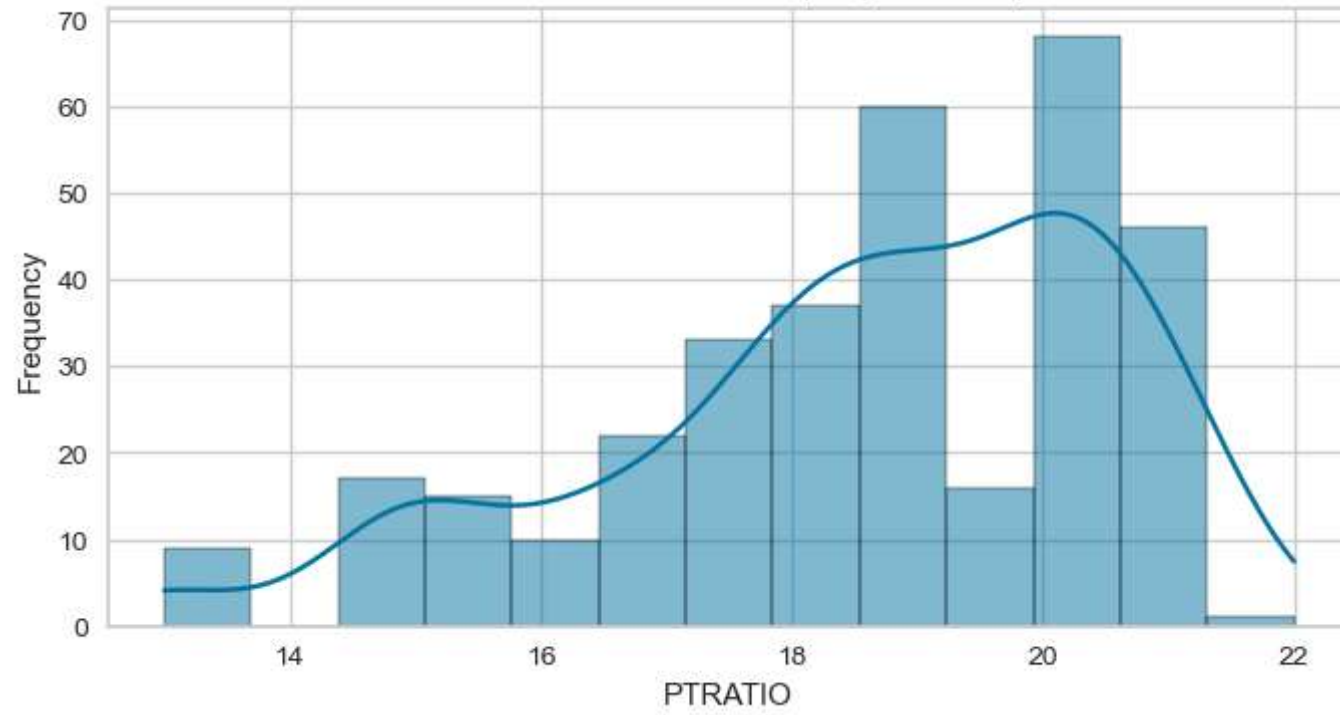


Distribution of RAD (Preprocessed)




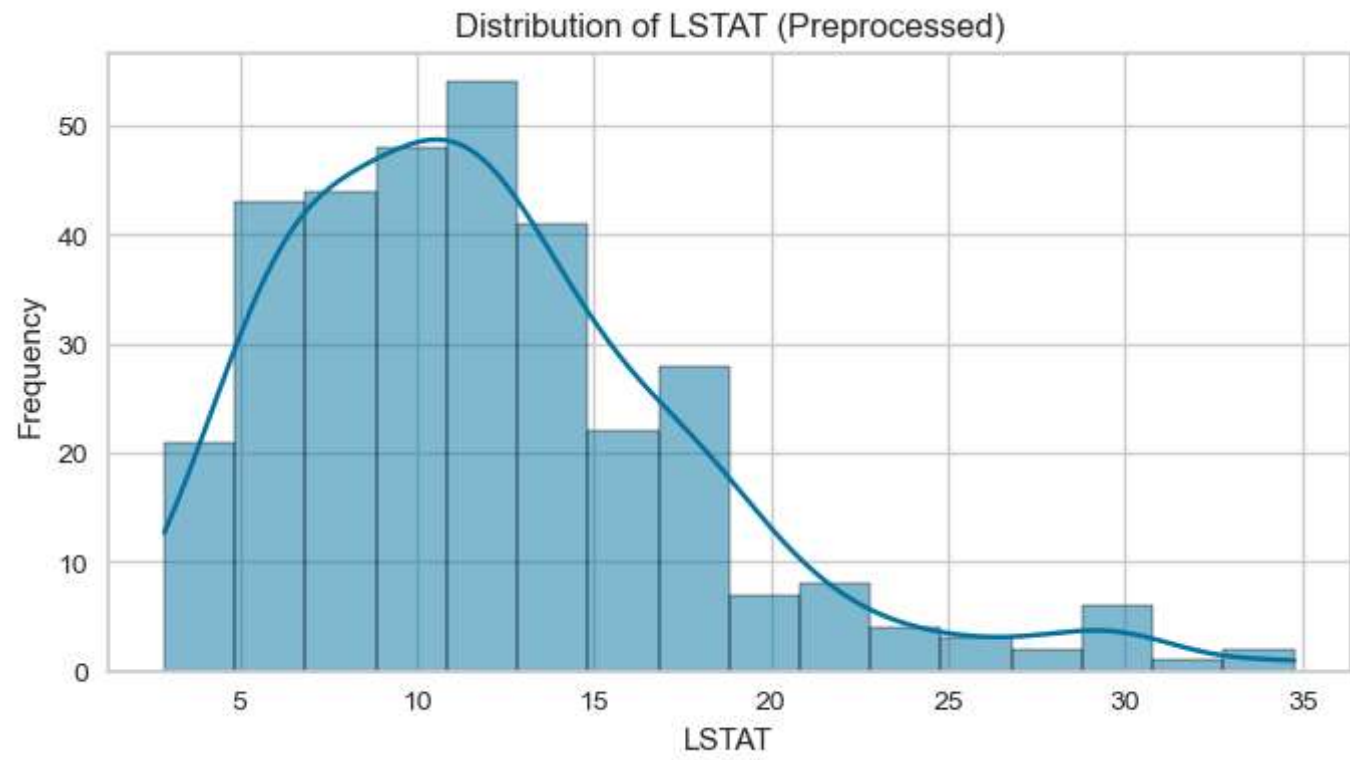


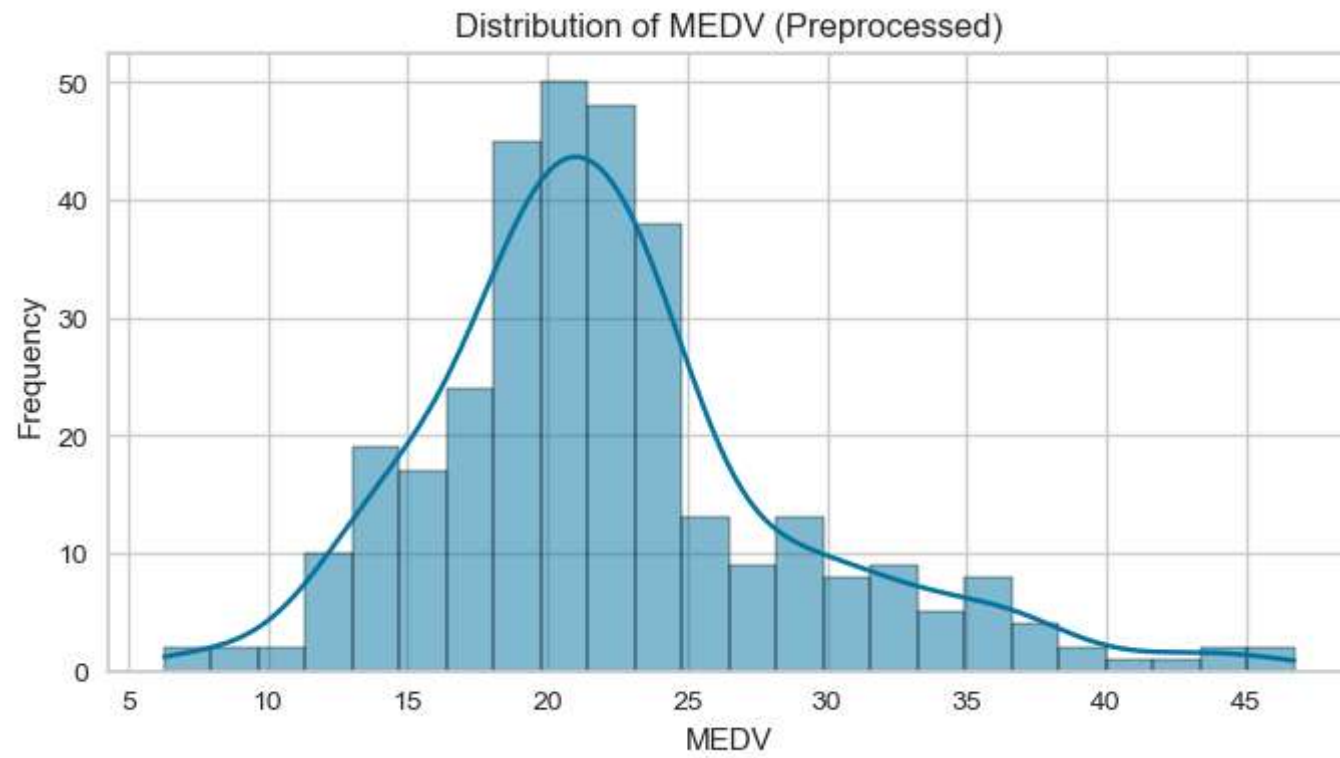
Distribution of PTRATIO (Preprocessed)



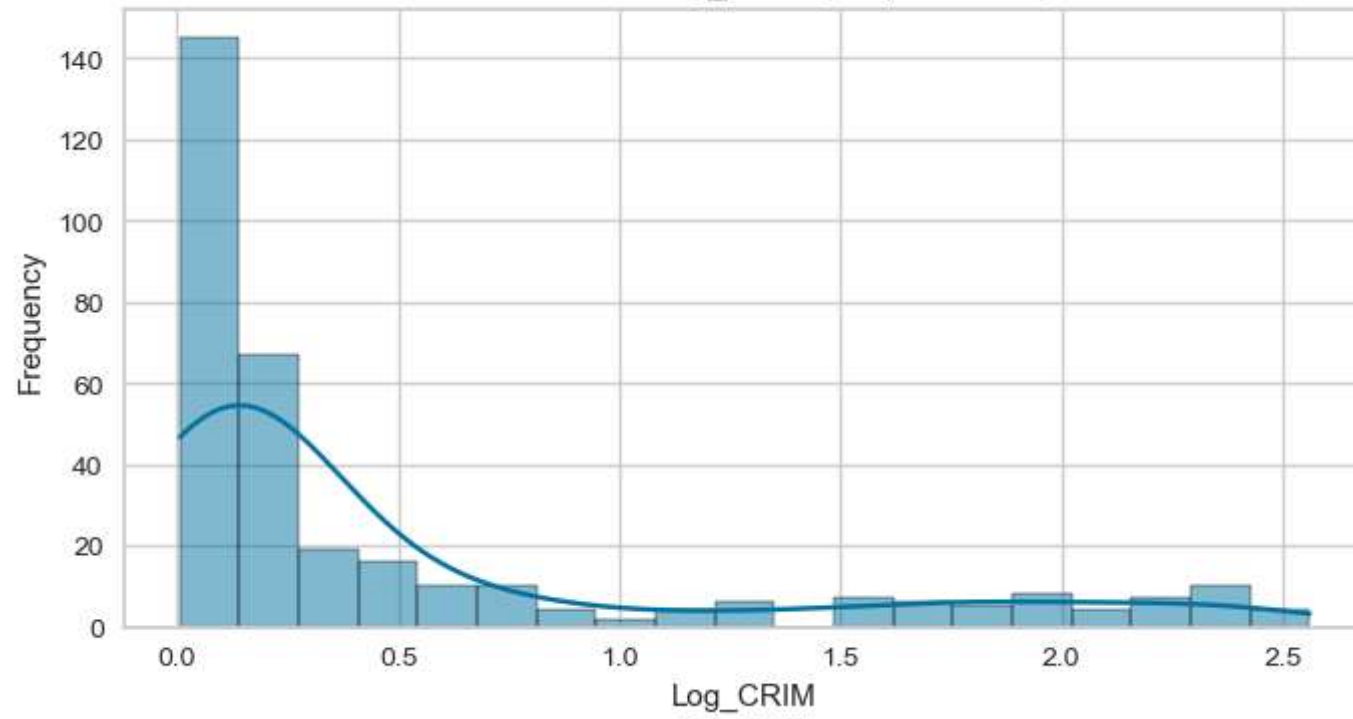
Distribution of B (Preprocessed)

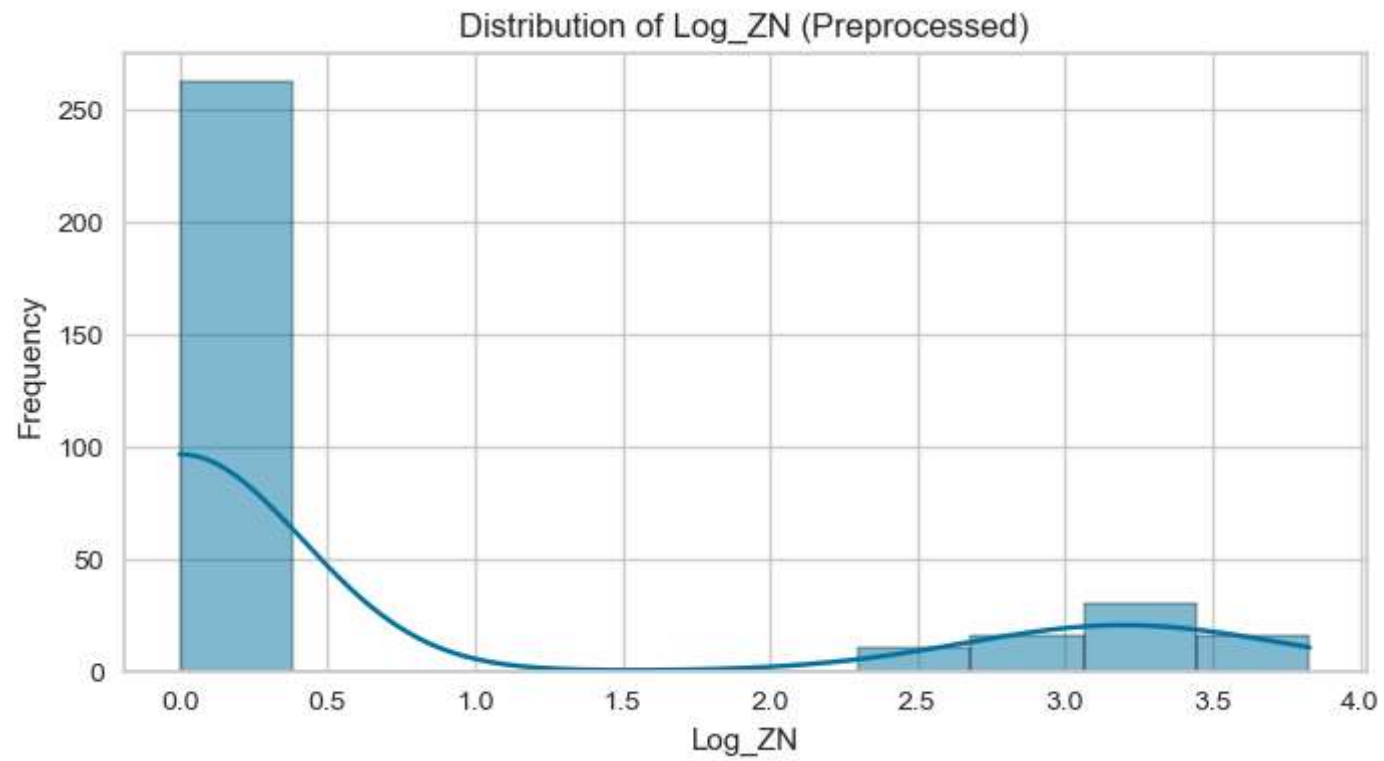


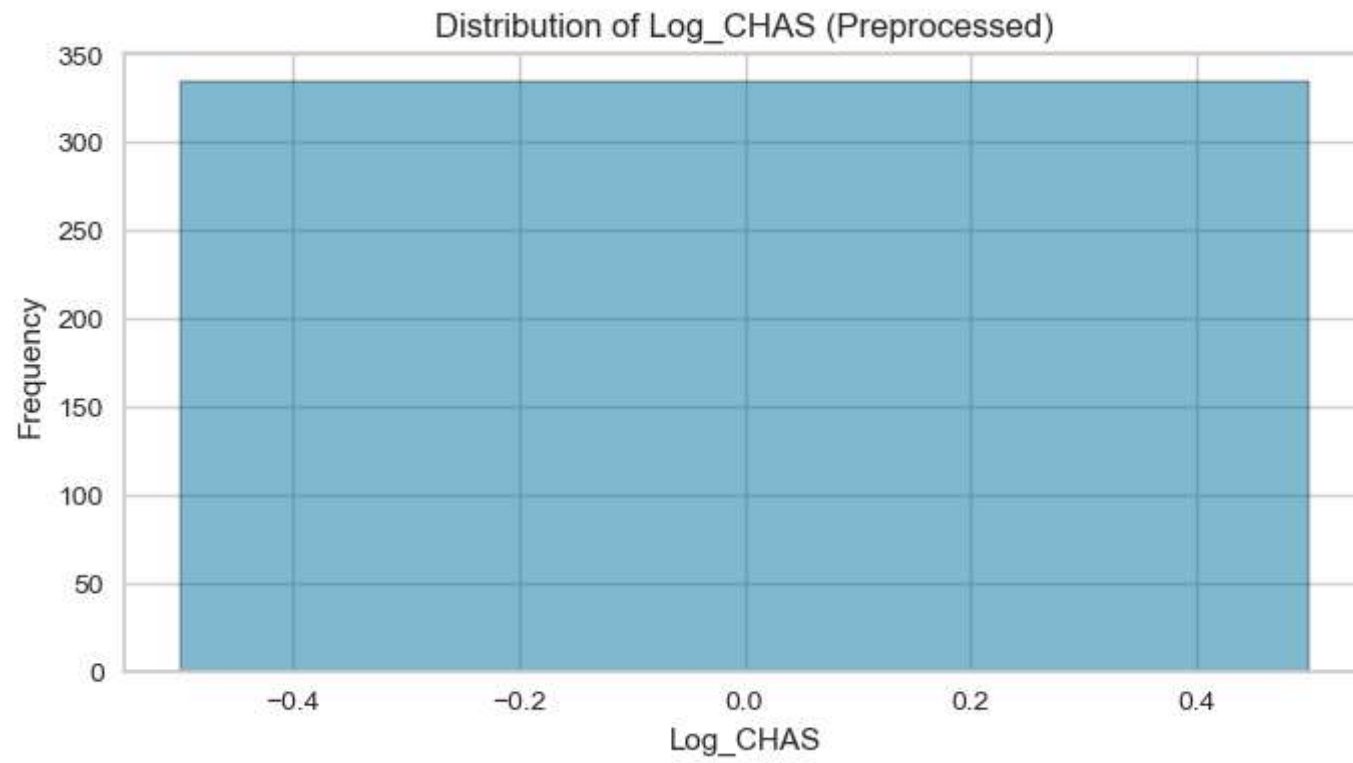


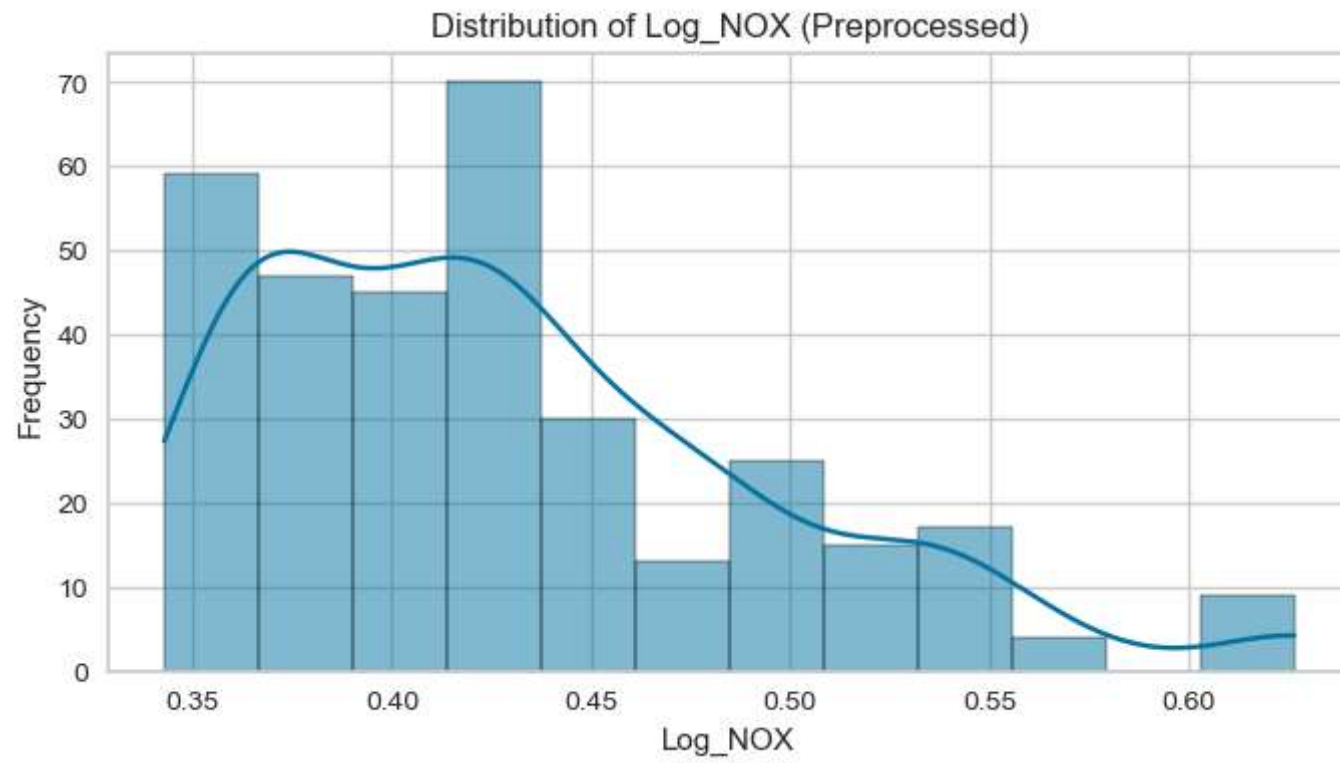



Distribution of Log_CRIM (Preprocessed)























[Step 2] Training and Evaluating Baseline Model

Baseline Model: Actual vs Predicted Values







```
C:\Users\vaibh\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
```




[Step 3] Training and Evaluating Main Model






```
C:\Users\vaibh\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
```

[Step 4] Performing Lasso Feature Importance Analysis



Best alpha: 0.47841076323135406
MSE: 13.048323531370661
R-squared: 0.68

==== Model Comparison Results =====

Baseline Model:

- MSE: 18.325
- R-squared: 0.815

Main Model:

- MSE: 9.213
- R-squared: 0.775

=====

